

## More on Transactions and Concurrency

---

---

---

---

---

---

---

## Serializability

- **Fallacy: being serializable is the same as being serial**
- Being serializable implies that the schedule is a correct schedule
  - It will leave the database in a consistent state
  - The interleaving is appropriate and will result in a state the equivalent of at least one schedule where the transactions are serially-executed.
  - It will achieve efficiency due to concurrent execution

---

---

---

---

---

---

---

## Checking for Serializability

$T_1$	$T_2$
$\text{read\_item}(X);$ $X := X + N;$	
	$\text{read\_item}(X);$ $X := X + M;$
$\text{write\_item}(X);$ $\text{read\_item}(Y);$	
	$\text{write\_item}(X);$
$Y := Y + N;$ $\text{write\_item}(Y);$	

- This schedule suffers from the lost update problem. It is not serializable as the final state of  $X = 7$  which is not possible with any sequential execution (either  $T_1:T_2$  or  $T_2:T_1$ ) of these transactions.

---

---

---

---

---

---

---

## Checking for Serializability (contd)

- **Serializability is hard to guarantee**
- Interleaving of operations happens at runtime through some scheduler
- Difficult to determine/predict how the operations of a schedule will be interleaved
- Practical Approach
  - Protocols that ensure serializability
    - Tradeoff (constrain how transactions are written, and their performance for serializability)
    - Reduce the problem of checking the whole schedule to checking only a committed projection of the schedule (I.

---

---

---

---

---

---

---

---

## Locks and Transactions

- Shared Locks (multiple outstanding locks)
  - Read Locks (many transactions can read at same time)
- Exclusive Locks (one transaction at a time)
  - Write Locks
- Rules (for transaction T)
  - T must obtain a read or write lock on X before reading X.
  - T must obtain a write lock before writing to X
  - T cannot obtain locks on items it already has locked
  - T must issue the unlock(X) operation after all read and write operations on X within T are finished.

---

---

---

---

---

---

---

---

## Using locks to guarantee serializability

- Initial values  
X=20, Y=30.
- Result of T1 → T2
  - X=50, X=80
- Result of T2 → T1
  - X=70, Y=50
- Does this work?
  - Is locking sufficient

(a)	T <sub>1</sub>	T <sub>2</sub>
	read_lock(Y);	read_lock(X);
	read_item(Y);	read_item(X);
	unlock(Y);	unlock(X);
	write_lock(X);	write_lock(Y);
	read_item(X);	read_item(Y);
	X := X + Y;	Y := X + Y;
	write_item(X);	write_item(Y);
	unlock(X);	unlock(Y);

---

---

---

---

---

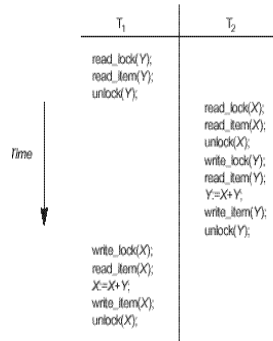
---

---

---

NO!

- Locking rules hold in this schedule
- Result of this schedule
  - $X=50, Y=50$
  - It does not match  $T1 \rightarrow T2$ , or  $T2 \rightarrow T1$
- So just locking does not suffice




---

---

---

---

---

---

---

---

## 2-phase locking

- Basic criteria (2 separable phases)
  - Locking Phase: Obtain all locks needed in the transaction [growing/expanding phase]
    - New locks may be obtained, none may be released.
  - Unlocking Phase: Release existing locks [shrinking phase]
    - Locks may be released, none may be obtained

---

---

---

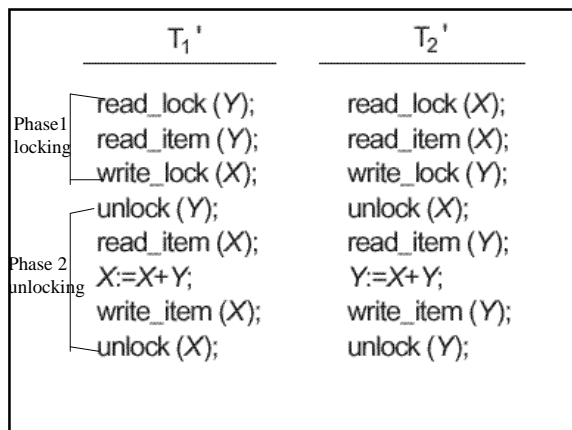
---

---

---

---

---




---

---

---

---

---

---


---

---

## All problems solved?

- Deadlock Problem

- T1 acquires Y, T2 acquires X. Both want the others lock and are unwilling to release their locks.




---

---

---

---

---

---

---

---

## 2-phase hierarchical locking

- Basic criteria (2 **separable** phases)

- Locking Phase: Obtain all locks needed in the transaction [growing/expanding phase]
  - New locks may be obtained, none may be released.
- Order locks, and obtain locks in predefined order.
  - Trade off deadlock avoidance for performance (less concurrency)
- Unlocking Phase: Release existing lock [shrinking phase]
  - Locks may be released, none may be obtained

---

---

---

---

---


---

---

---

## Solved? Yes!

T1	T2
Read_lock(y)	Write_lock(Y)
Read_item(y)	Read_lock(X)
Write_lock(x)	Read_Item(X)
Unlock(y)	Unlock (x)
Unlock(X)	Unlock (y)



With the above transactions, schedule on right can never happen.  
Unlock can happen in any order.

---

---

---

---

---

---

---

---

## Conflict Equivalence

- Two schedules are said to be conflict equivalent if the order of any two **conflicting operations** is the same in both schedules
- Conflicting Operations (different transactions)
  - Read after Write (RAW)
    - $W_1(X), R_2(X)$
  - Write after Read (WAR)
    - $R_2(Y), W_1(Y)$
  - Write after Write (WAW)
    - $W_1(Y), W_2(Y)$

---

---

---

---


---

---

---

---

## Conflict Serializability




---

---

---

---

---

---

---

---

## Testing Conflict Serializability

- Create a precedence graph.
  - Create a node for each transaction.
  - Create a dependency (line from one node to another) for every conflict.
- Test for cycles in precedence graph.
- No cycles  $\Rightarrow$  conflict serializable
- Cycles  $\Rightarrow$  no serial schedule exists that is conflict equivalent to original schedule.

---

---

---


---

---

---

---

---




---

---

---

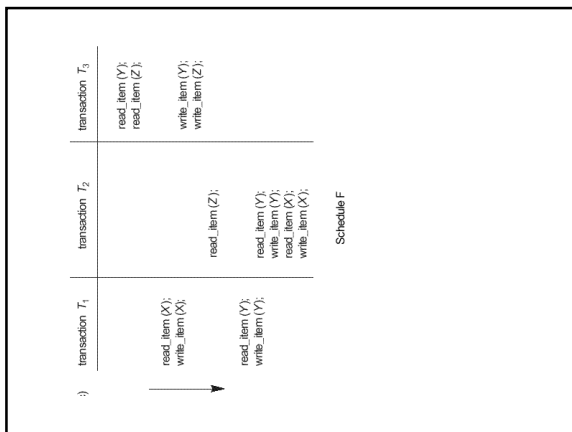
---

---

---

---

---




---

---

---

---

---

---

---

---

## View Serializable

- Slightly weaker notion of serializability when compared to conflict-serializability.
- Premise
  - Each read operation of a transaction reads the result of the same write operation in both schedules
  - The write operation of each transaction must produce same results

---

---

---

---

---

---

---

---

## Relationship between View and Conflict Serializability

- The two are the same under the “constrained write assumption” which assumes any transaction T that writes a value X (in other words **no blind writes**)
  - Reads OLD VALUE OF X
  - New X = F(OLD VALUE OF X)
- Example
  - T1: r1(X), w1(X)
  - T2: w2(X)
  - T3: w3(X)
  - The schedule r1(X);w2(X);w1(X);w3(X); is view equivalent but not conflict serializable

---

---

---

---

---

---

---